Response to reviewer comments on "Generalized Jeffreys's approximate objective Bayes factor: Model-selection consistency, finite-sample accuracy, and the Bayesian evidence in early-phase clinical trials"

Puneet Velidi, Zhengxiao Wei, Evan Strasdin, Shreena Nisha Kalaria, Yimeng Liu, Céline M. Laumont, Brad H. Nelson, Farouk S. Nathoo

October 30, 2026

Response to Associate Editor's Comments

The paper introduces a generalized form of Jeffreys's approximate objective Bayes factor (eJAB). It was read by myself and two expert referees. The authors establish consistency and assess finite-sample accuracy via simulations. There were numerous major concerns identified by the reviewers:

1. *'more exploration of the connection to test-based BF (Johnson, 2005)'*

   We originally selected test-statistic Bayes factors (TSBFs) for two main reasons: their computational convenience, as they admit closed-form expressions (although TSBFs are not 'approximate' Bayes factors in the same sense as eJAB), and the lack of readily available MCMC-based Bayes factors for chi-squared and other nonparametric tests considered in our simulations. In our response to Reviewer 1's comment, we now discuss the assumptions and asymptotic properties of three versions of the TSBF and compare them with those of eJAB. We have also revised the simulation studies to avoid comparisons between eJAB and the minimum Bayes factor version of TSBF.

2. *'significant concerns about the regularity condition R2'*

   In our response to Reviewer 1, we have carefully addressed concerns regarding the regularity condition R2 required for model-selection consistency. First, we have discussed asymptotic properties of TSBF in relation to R2 in our first response to Reviewer 1. In our second response to Reviewer 1, we have made comparisons with the BIC approximation to the Bayes factor and the conditions required for model selection consistency in relation R2. Finally, we have conducted an empirical investigation of R2 based on data obtained from four large scale replication studies, where we observe realizations of D_n under H_1 and comment on the behavior of D_n with increasing sample size.

3. *'significant concerns about the analysis in Section 4.'*

   We have revised the analysis in response to Reviewer 1's suggestions. The application is now restricted to 16,563 primary- and secondary-outcome analyses from 1,790 early-phase trials, with detailed inclusion

and exclusion criteria provided in Appendix C. One-sided tests are now considered explicitly in addition to two-sided tests. In the revision we have defined a one-sided version of eJAB and discussed its asymptotic properties.

4. *Addition of new coauthor:* We have added one new coauthor to the manuscript. Evan Strasdin is a graduate student who conducted an empirical study that forms the second part of our response to the second comment of Reviewer 1. We have added this study as a new section in the revision.

5. *Minor modification to the title of the manuscript:* We have changed the title of the manuscript from '*Generalized Jeffreys's approximate objective Bayes factor: Model-selection consistency, finite-sample accuracy, and statistical evidence in 71,126 clinical trial findings*' to '*Generalized Jeffreys's approximate objective Bayes factor: Model-selection consistency, finite-sample accuracy, and the Bayesian evidence in early-phase clinical trials*'. This change is required as part of our response to the fifth comment of Reviewer 1, who suggested that we only look at early phase trials (phase I and II) in our analysis.

Response to Reviewer 1

- *The proposed BF depends on the test-statistic of a chi-squared test with dimension q. It also depends on sample size n. The connection to the test-based BF (Johnson, 2005) should be explored in more detail, both theoretically and through simulation. For example, Johnson also considers BFs based on the F- (rather than the chi-squared) statistic, which also depend on sample size. In fact, you do use F-tests in simulation items 4 and 5. A closer theoretical comparison with BFs based on a chi-squared test statistic is warranted. The derivation in Johnson is based on certain asymptotic assumptions (local alternatives); how do these compare with assumptions R1 and, in particular, R2?*

Thank you for raising this question. We have expanded the discussion to distinguish three related constructions: the local alternative prior, the maximum marginal likelihood estimate (MMLE), and the nonlocal alternative prior. Johnson's TSBF models the sampling distribution of a classical test statistic rather than the full-data likelihood. For the chi-squared test, the statistic has a central chi-squared distribution under $\mathcal{H}_0$ and a noncentral chi-squared distribution under $\mathcal{H}_1$, yielding a closed-form 'Bayes factor' that depends on the statistic x_n, its degrees of freedom q, the sample size n, and an alternative-model dispersion parameter $\varphi = (c+1)/n$, where $c = \mathcal{O}(n)$ denotes the Dirichlet concentration parameter,

$$\mathcal{H}_0 : \; x_n \sim \chi^2_q, \tag{1}$$

$$\mathcal{H}_1 : \; x_n \mid \gamma \sim \chi^2_q(\gamma), \qquad \gamma \mid \varphi \sim \text{Gamma}\left(\frac{q}{2}, \frac{\varphi}{2}\right), \tag{2}$$

$$TSBF_{10} = \left(\frac{\varphi}{\varphi+1}\right)^{\frac{q}{2}} \exp\left\{\frac{x_n}{2(\varphi+1)}\right\}. \tag{3}$$

Johnson (2005) considers maximizing the marginal density of the observed test statistic under $\mathcal{H}_1$ over the dispersion parameter. This produces the minimum BF_{01} within the specified class (Held and Ott, 2018). It is therefore an upper bound and the least conservative assessment of evidence against $\mathcal{H}_0$. Our previous manuscript compared eJAB only with this MMLE version, rather than with the broader class of TSBFs (Johnson, 2005, 2008; Yuan and Johnson, 2008; Johnson, 2013). Because this quantity is maximized in favor of $\mathcal{H}_1$, it is bounded below by one and equals one when the observed statistic does not exceed its expectation under $\mathcal{H}_0$, producing the characteristic 'hockey-stick' behavior. We agree that comparing an approximate Bayes factor with a minimum Bayes factor is an apples to oranges comparison. Accordingly, we have removed the comparisons with the MMLE-TSBF and replaced them with comparisons to the TSBF more generally.

The asymptotic assumptions differ in their intent. To test the hypotheses $\mathcal{H}_0 : \boldsymbol{\theta} = \boldsymbol{\theta}_0$ versus $\mathcal{H}_1 : \boldsymbol{\theta} \neq \boldsymbol{\theta}_0$, Johnson's original chi-squared derivation assumes local alternatives satisfying $\boldsymbol{\theta} - \boldsymbol{\theta}_0 = \mathcal{O}_p(n^{-\frac{1}{2}})$. In other words, a local alternative not only centers a unimodal, symmetric prior distribution at the null value but also specifies that the prior scale shrinks with the sample size. One consequence of this assumption is the existence of a nondegenerate asymptotic distribution for the test statistic under $\mathcal{H}_1$—specifically, a noncentral chi-squared limit with an $\mathcal{O}_p(1)$ noncentrality parameter (Johnson, 2005, Lemma 1). It is not a condition that ensures model-selection consistency. Under this local-alternative regime, regularity condition R2 will not hold, and model-selection consistency of eJAB is not guaranteed. However, we note that the TSBF is also not model-selection consistent under the same assumption. Under $\mathcal{H}_0$, $x_n \xrightarrow{D} \chi^2_q$ and $\varphi \to \varphi_0 > 0$ as $n \to \infty$, so $TSBF_{10}$ in Eq. 3 is nondegenerate and almost surely positive. Thus, under Johnson's local-alternative framework, the asymptotic TSBF exists in a nondegenerate form, but neither eJAB nor the TSBF is model-selection consistent. By contrast, eJAB is based on three asymptotic assumptions: (1) Asymptotic normality of the posterior distribution; (2) R1: $p \xrightarrow{D} \text{Unif}(0,1)$ as $n \to \infty$ under $\mathcal{H}_0$; (3) R2: $D_n = -\sqrt{n} \cdot p \ln p \xrightarrow{P} 0$ as $n \to \infty$ under $\mathcal{H}_1$. The first assumption along with the assumed prior gives us the form for eJAB. The latter two assumptions are sufficient for model-selection consistency (i.e., Theorem 2.3). Under these assumptions, the same argument shows that $TSBF_{10}$ in Eq. 3 does not converge to zero under $\mathcal{H}_0$ and remains

model-selection inconsistent. We have revised the manuscript to clarify the differences in assumptions and their implications.

We further consider the model selection consistency of TSBF in subsequent work, across three recent papers, the Bayes factor function (BFF) retains the same core structure—a family of TSBFs evaluated over alternative hypotheses indexed by standardized effect size ω—while the principal extensions concern the nonlocal priors used to define those alternatives, expanding the first-order moment and gamma priors to arbitrary-order moment and inverse-moment/inverse-gamma families and examining how their shape and dispersion affect evidence accumulation, operating characteristics, and aggregation across replicated studies (Johnson et al., 2023; Datta et al., 2025a,b). The following expressions give the BFF for the chi-squared test proposed by Johnson et al. (2023),

$$\mathcal{H}_0 : \; x_n \sim \chi^2_q, \tag{1}$$

$$\mathcal{H}_1 : \; x_n \mid \gamma \sim \chi^2_q(\gamma), \qquad \gamma \mid \tau^2 \sim \text{Gamma}\left(\frac{q}{2}+1,\, \frac{1}{2\tau^2}\right), \tag{4}$$

$$BFF_{10}\left(x_n \mid \tau^2\right) = (1+\tau^2)^{-\frac{q}{2}-1}\left(1+\frac{\tau^2 x_n}{q(1+\tau^2)}\right)\exp\left\{\frac{\tau^2 x_n}{2(1+\tau^2)}\right\}, \tag{5}$$

$$\tau^2 = \frac{n}{q}\omega^2. \tag{6}$$

They assume the alternative in Eq. 4 so that the prior density on the noncentrality parameter vanishes at the origin. Since $\text{mode}(\gamma) = q\tau^2$, Eq. 6 matches the prior mode to the noncentrality $n\omega^2$ corresponding to a specified standardized effect size ω. This first-order nonlocal prior yields faster accumulation of evidence in favor of a true null than the corresponding local prior. Under $\mathcal{H}_0$, $x_n = \mathcal{O}_p(1)$ and, for every fixed $\omega > 0$, $\tau^2 = \mathcal{O}(n)$, so $BFF_{10} = \mathcal{O}_p\left(n^{-\frac{q}{2}-1}\right) \xrightarrow{P} 0$ as $n \to \infty$. The faster polynomial accumulation of null evidence motivates the higher-order nonlocal priors considered in Datta et al. (2025a). Under a fixed alternative with true standardized effect $\omega_0 > 0$, the true noncentrality parameter satisfies $\gamma = n\omega_0^2$, and hence $x_n = n\omega_0^2 + \mathcal{O}_p(\sqrt{n})$. For any fixed nonzero BFF ordinate $\omega > 0$, $\tau^2 = \mathcal{O}(n)$, and therefore the exponential term in Eq. 5 dominates its polynomial penalty. In particular, $\frac{1}{n}\ln BFF_{10} \xrightarrow{P} \frac{\omega_0^2}{2}$, so $BFF_{10} \xrightarrow{P} \infty$ as $n \to \infty$. Thus, for every fixed nonzero effect-size ordinate ω, the BFF is pointwise model-selection consistent.

- *More theoretical motivation and discussion are needed for regularity condition R2. It is a stronger condition than assuming that the traditional minimum BF calibration, $-e \cdot p \ln p$ (Sellke et al., 2001), converges to zero under $\mathcal{H}_1$ as n goes to infinity. So how strong is condition 2? Is this assumption justifiable?*

Assume that, under $\mathcal{H}_1$, the p-value is bounded away from 1 for all sufficiently large n. Then, $-e \cdot p \ln p \xrightarrow{P} 0$ is equivalent to $p \xrightarrow{P} 0$ because $\lim_{x \to 0^+} x \ln x = 0$. Thus, requiring the traditional minimum-BF calibration to converge to zero imposes no restriction on the rate at which the p-value converges to zero under $\mathcal{H}_1$. With regard to the strength of assumption R2: $D_n = -\sqrt{n} \cdot p \ln p \xrightarrow{P} 0$ as $n \to \infty$ under $\mathcal{H}_1$, we first note that it is not sufficiently strong to ensure model-selection consistency of the BIC approximation to the Bayes factor, as formulated in Eq. 7. Let $Q_{\chi^2_q}(1-p)$ denote the quantile function of a chi-squared distribution with q degrees of freedom, evaluated at tail probability $1-p$, where $p \to 0^+$. Suppose that the p-value is computed from a likelihood-ratio test so that, under standard regularity conditions,

$$BF_{01}^{(\text{BIC})} = n^{\frac{q}{2}} \exp\left\{-\frac{1}{2}Q_{\chi^2_q}(1-p)\right\}, \tag{7}$$

and that R2 holds, but $-n^a \cdot p \ln p \xrightarrow{P} 0$ for $a > \frac{1}{2}$ as $n \to \infty$ under $\mathcal{H}_1$. Then, $\exists \epsilon > 0$ such that $\mathbb{P}(-n^a \cdot p \ln p < \epsilon) \nrightarrow 1$, which, in turn, implies that $\mathbb{P}(BF_{01}^{(\mathrm{BIC})} < \epsilon) \nrightarrow 1$ for $q > 1$ under $\mathcal{H}_1$.

More specifically, using the asymptotic expansion of the incomplete gamma function (e.g., Daouia et al., 2025, Lemma A.2), $\frac{1}{2} Q_{\chi^2_q}(1-p) = -\ln p + (\frac{q}{2} - 1) \ln(-\ln p) - \ln \Gamma(\frac{q}{2}) + o(1)$. It follows that $BF_{01}^{(\mathrm{BIC})} \approx \Gamma(\frac{q}{2})\, n^{\frac{q}{2}} \cdot p\, (-\ln p)^{1-\frac{q}{2}}$, which requires a substantially stronger condition than R2 for model-selection consistency when $q > 1$.

We next consider condition R2 from a more practical standpoint. We conduct a study examining the realized values of D_n in four large-scale replication projects. Each project selected a collection of previously published findings and attempted to replicate them under a preregistered protocol, so that each original result was paired with an independently conducted replication of the same effect.

- **EERP**, the Experimental Economics Replication Project (Camerer et al., 2016), which attempted to replicate laboratory experiments published in the *American Economic Review* and the *Quarterly Journal of Economics* between 2011 and 2014;
- **RPP**, the Reproducibility Project: Psychology (Open Science Collaboration, 2015), which attempted to replicate a collection of findings drawn from three high-impact psychology journals;
- **RPCB**, the Reproducibility Project: Cancer Biology (Errington et al., 2014), which considered experiments reported in high-impact preclinical cancer biology papers; we use the project's final analysis (Errington et al., 2021); and
- **SCORE**, Systematizing Confidence in Open Research and Evidence (Tyner et al., 2025), which assembled and considered replication of a large corpus of published claims from the social and behavioral sciences.

For each result in the four projects, we extracted the original p-value and sample size n, and the p-value and sample size of each of its replication attempts. Each result has at least one and possibly several replication attempts. Our objective is to examine realized values of D_n under H_1 as a function of sample size. To do so, we create an $\mathcal{H}_1$ subset corresponding to replicated results where it is safe to assume that $\mathcal{H}_1$ is true. A result is included in the $\mathcal{H}_1$ subset when its original p-value satisfies $p \leq 0.005$ and all of its replication attempts also reached $p \leq 0.005$. Given that the effect was found at this level in the original study and in every replication, we make the assumption that $\mathcal{H}_1$ is true for these results. This gives realized values of D_n under $\mathcal{H}_1$: one from each original study and one from each replication attempt. Table 1 reports the counts of the results reported and used in our study.

Two kinds of recorded p-value were treated separately. Where a p-value was recorded as 0, we computed $\ln p$ from the reported test statistic and, where the test has them, its degrees of freedom. Of the plotted values, 51 were obtained this way, including the value at the largest sample size; 12 values (six SCORE original results and six SCORE replication attempts) could not be computed and were dropped. Where a p-value was reported as an upper bound, such as $p < .001$, rather than as a value, we removed it. Such left-censored p-values occur only in SCORE. SCORE labels its original p-values; we removed the 37 it labels as less-than or threshold values, and the one with no label. SCORE does not label replication p-values, so we treated a replication p-value recorded as an exact power of ten, such as .001, as censored unless we were able to recompute the p-value directly from the reported test statistic. This removed two replication values where the p-value is left censored and cannot be recalculated from the reported test statistic.

Figure 1 plots the realized values of D_n against the sample size n, which spans roughly seven orders of magnitude. A clear pattern emerges indicating that the distribution of D_n is shifted more toward zero as n increases. The value of D_n at the largest value of n, $n = 49{,}196{,}731$ with $D_n = 4.09 \times 10^{-29}$, is circled in the figure. Overall, the study suggests that R2 is practically feasible.

The two largest sample sizes depicted in Figure 1 are one SCORE claim and its replication. The original study from 2014 (Travers and Krezmien, 2018) examined autism identification among US public school students, using counts of students with autism from the IDEA Data Center and enrolment from the NCES Digest of Education Statistics. The claim SCORE tested is that, compared with California, Ohio has lower odds of autism identification, assessed by a logistic regression of autism identification on state, with each student as one observation. The replication repeated this analysis with the same two sources for 2017, on 49,196,731 students in 46 states; SCORE records 7,977,660 students for the original. These sample sizes are counts of enrolled students taken from administrative records, not

Project	Field	Results				(5) Plotted values	
		(1) released	(2) with p, n	(3) $p \leq 0.005$	(4) $\mathcal{H}_1$ subset	original	replication
EERP	Experimental economics	18	18	7	5	5	5
RPP	Psychology	167	143	67	16	16	16
RPCB	Preclinical cancer biology	158	115	61	15	15	21
SCORE	Social and behavioral sciences	331	311	184	102	58	107
	Total	674	587	319	138	94	149

Table 1: Counts at each step of the selection, by project. Column 1 *Released*: the number of results in each project's released data. A result is a study for EERP (18) and RPP (167, of which 100 were replicated), an effect for RPCB (158 effects from 23 papers, given in 188 rows), and a claim for SCORE (331 claims with at least one replication report). Column 2 *With p and n*: the results in Column 1 for which both the original p-value and the original sample size are recorded. Column 3 *Original $p \leq 0.005$*: the results in Column 2 whose original p-value satisfies $p \leq 0.005$. Column 4 $\mathcal{H}_1$ *subset*: the results in Column 3 for which every replication attempt reached $p \leq 0.005$. A result with no recorded replication p-value is not included; this excludes 20 RPP results and one RPCB result. We also excluded RPP study 69. The RPP data record $p = 10^{-6}$ for both its original and its replication, with no test statistic from which either value could be checked, whereas the reported p-values in the same data are $p < .05$ for the original and $p = .05$ for the replication. For SCORE, a replication attempt is one replication report. Columns 5 and 6 *Plotted values*: the number of values of D_n shown in Figure 1, one for each original result and one for each replication attempt in the $\mathcal{H}_1$ subset. This excludes left censored p-values that cannot be recomputed from the reported test statistics. For RPCB and SCORE the replication column exceeds the $\mathcal{H}_1$ subset because a result can have more than one replication attempt (up to three in RPCB and four in SCORE), each contributing its own value.

participants recruited for a study. The estimated effects are small, with odds ratios of 0.935 in the original and 0.912 in the replication ($z = -6.31$ and $z = -12.32$).

This study is reported in Appendix E of the Supplementary Material, and the data, code and counts for Table 1 and Figure 1 are in the GitHub repository. We refer to it at the end of Section 3.1 of the revision.

- *The analysis described in Section 4 is unconvincing for several reasons. First, clinical trial findings are usually based on one-sided tests at significance level $\alpha = 0.025$.*

In the revised analysis, we determine the test direction from the reported statistical method. For explicitly one-sided scalar tests ($q = 1$), we apply the one-sided eJAB discussed below; for two-sided or otherwise unrestricted tests, we apply the two-sided eJAB. This classification yields 427 one-sided and 16,136 two-sided or unrestricted results in the revised analysis. When the test direction is not reported, we assign the analysis to the latter category. Supplementary Material Appendix C.5 provides the classification rules and describes the treatment of reported results, including censored p-values. For the Bayes/NHST contradiction proportions in Figure 3, a contradiction requires eJAB > 1 (one-sided or two-sided as appropriate) and $p \leq .05$ for two-sided tests or $p_1 \leq .025$ for one-sided tests.

Equality is included consistently in the text and executable analysis. Comparing numerical reports as displayed adds 65 results at their direction-specific threshold, including 32 contradictions, for totals of 6,474 significant results and 493 contradictions. Reported bounds retain their original inequality signs; rounding of numerical reports is acknowledged in Appendix C.5.

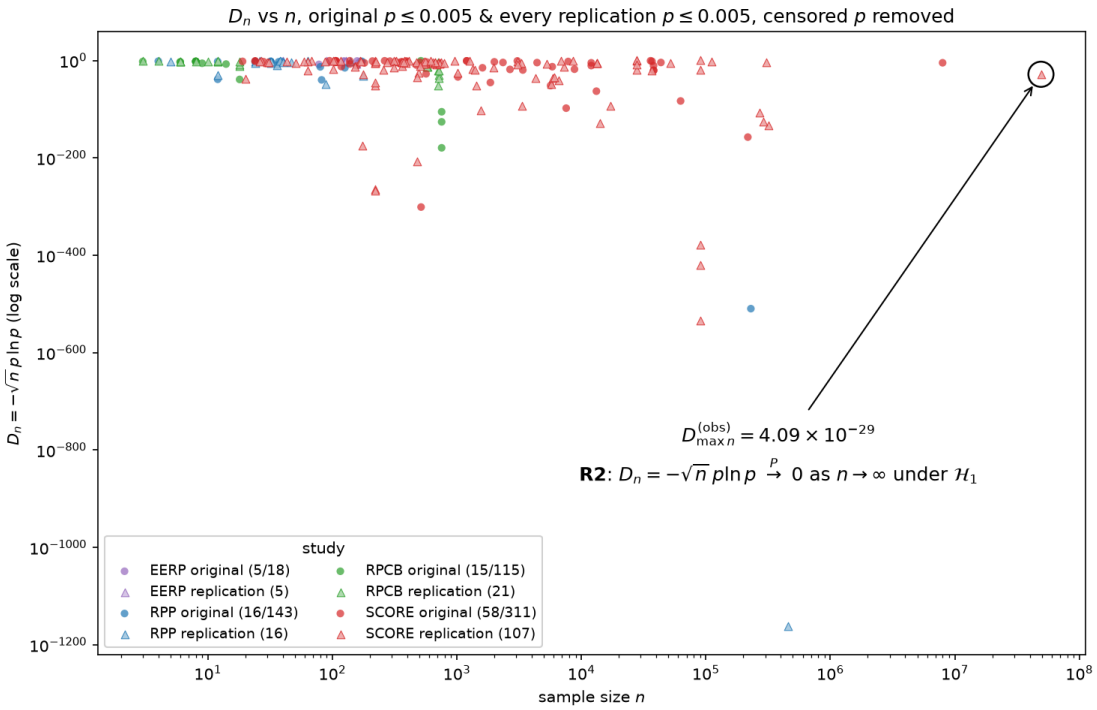


Figure 1: Realized values of $D_n = -\sqrt{n} \cdot p \ln p$, on a logarithmic scale, plotted against sample size n for the $\mathcal{H}_1$ subsets of the four replication projects (original $p \leq 0.005$ and every replication attempt $p \leq 0.005$). Circles denote original results and triangles replication attempts, with symbols colored by project. Legend entries give the number of plotted values; for original results this is shown relative to column 2 of Table 1. p-values reported as bounds are excluded. The circled value, $D_{\max n}^{(\text{obs})} = 4.09 \times 10^{-29}$, is the observed D_n at the largest sample size.

- *The proposed BF, however, is intrinsically two-sided, as it is based on a chi-squared statistic. Can you amend eJAB to the one-sided setting, taking into account the direction of the effect? A one-sided Bayes factor may need different priors and assumptions from those used here.*

  We have extended eJAB to the one-sided setting, e.g., $\mathcal{H}_0 : \theta = \theta_0$ versus $\mathcal{H}_+ : \theta > \theta_0$, and added the corresponding discussion, which explicitly incorporates the direction of the effect and the assumptions required for the one-sided Bayes factor. We have accordingly provided Appendix C.5 in the Supplementary Material.

  The encompassing-prior correction of Wagenmakers (2022, pp. 20–22) multiplies the Bayes factor for the null against an unrestricted alternative by the ratio of prior to posterior probability assigned to the specified direction. For a scalar right-sided test, we use a symmetric local prior centered at θ_0, assigning prior mass $1/2$ to the right-sided alternative and approximating the corresponding posterior mass by $1 - p_1$. With $p = 2\min\{p_1, 1 - p_1\}$, combining the null-centered Savage–Dickey approximation with the directional correction yields

$$eJAB_{0+} = \frac{\sqrt{n}}{2(1-p_1)} \exp\left\{-\frac{1}{2} Q_{\chi^2_1}(1-p)\right\}.$$

  When $q = 1$, the two $eJAB$ definitions are related by

$$eJAB_{0+} = \frac{n^{\frac{1}{2(1-n)}}}{2(1-p_1)} eJAB_{01}^{\frac{n}{n-1}}.$$

Under $\mathcal{H}_0$, **R1** implies $eJAB_{01} \xrightarrow{P} \infty$. Since $n^{\frac{1}{2(1-n)}} \to 1$ and $(2(1-p_1))^{-1} \geqslant 1/2$, we have $eJAB_{0+} \xrightarrow{P} \infty$. Under $\mathcal{H}_+$, **R2** implies $eJAB_{01} \xrightarrow{P} 0$. Assuming additionally **R3**: $(1-p_1)^{-1} = \mathcal{O}_p(1)$ under $\mathcal{H}_+$, it follows that $eJAB_{0+} \xrightarrow{P} 0$. Therefore, the one-sided $eJAB$ is also model-selection consistent.

- *Second, if you are concerned that 'failures at later phases may reflect type I errors, where the null hypothesis was incorrectly rejected at an earlier phase', then you should investigate only early-phase studies (Phases I and II) in your analysis. However, the analysis includes Phase II and even Phase IV studies.*

  To address this concern, we now restrict the manuscript primary- and secondary-outcome analysis to 16,563 eligible analyses from 1,790 exact Early Phase I, Phase I, and Phase II interventional trials. Combined-phase trials and Phase III and IV trials are now excluded. We report the Bayes/NHST contradiction-rate profiles and posterior densities of contradiction rates separately for the three classifications so that Early Phase I and Phase I are not combined. The observed contradiction proportions among direction-specifically significant results are 6/94 (6.38%), 28/927 (3.02%), and 459/5,453 (8.42%), respectively. These differ from the proportions among all eligible analyses in Figure 2 and the posterior summaries in Figure 3.

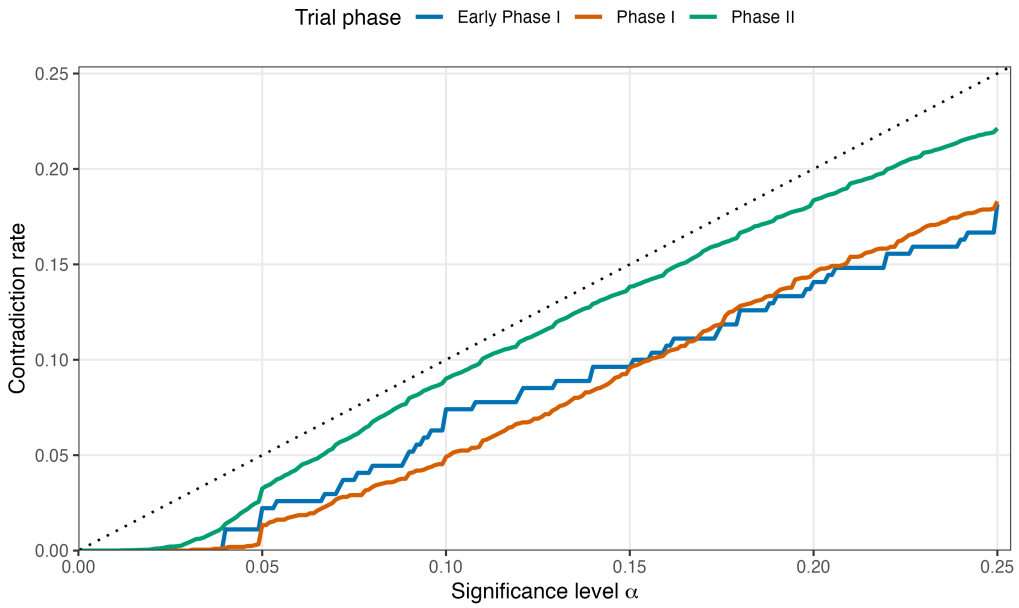


Figure 2: Bayes/NHST contradiction rate versus significance level among all eligible analyses in each trial phase. The sweep uses p for two-sided/unrestricted tests and $2\min\{p_1, 1-p_1\}$ for one-sided tests. The latter also admits opposite-tail results and differs from the directional rule used for the reported counts.

- *Third, to justify such broad claims about type I errors in clinical trial findings, this study should be much more rigid, with a predefined research question and a study protocol appropriately preregistered before the analysis is conducted. But the study is purely exploratory, and it remains unclear how the 71,126 CTGs were selected, why the described analysis pipeline ('We fit a ... model', etc.) is appropriate, and so on. On the positive side, I appreciate the availability of code on OSF.*

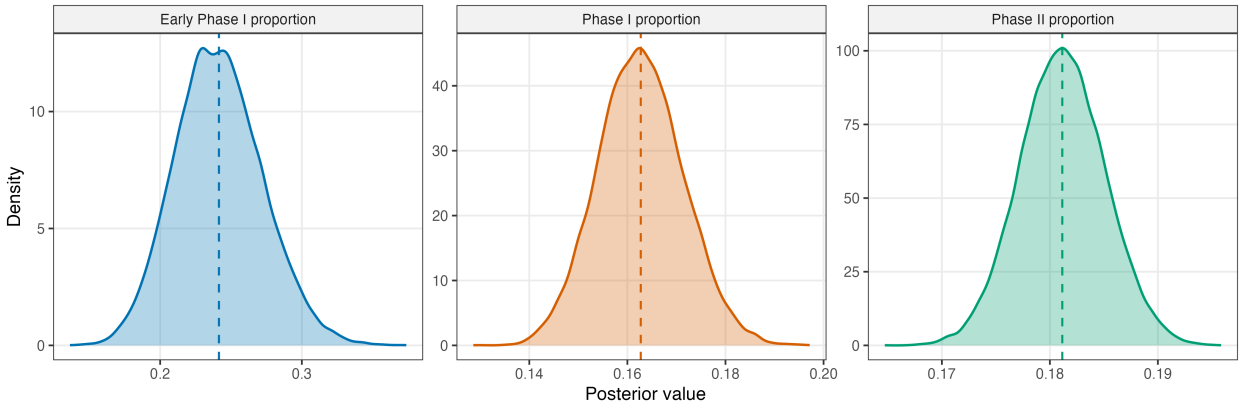


Figure 3: Study-level posterior distributions of the contradiction proportion among direction-specifically significant results, shown separately for the three trial phases. Dashed lines mark posterior means. The Dirichlet$(1, 1, 1, 1)$ prior applied to each study affects these estimates, which differ from the raw observed proportions.

Our analysis was not prospectively preregistered, and we therefore make more cautious statements about Bayes/NHST contradictions in the revision. Our analysis is certainly exploratory, the goal being to reexamine published clinical trial results under a Bayesian lens when only limited information and the original data are not available.

We provide detailed inclusion and exclusion criteria for the clinical trials in Supplementary Material Appendix C, together with the sequential flow from 279,927 analyses in 22,454 studies to the final 16,563 analyses in 1,790 trials. Both significant and nonsignificant eligible results are retained. The appendix specifies the test mappings, sample-size and dimension rules, and treatment of reported inequalities. For left-censored values such as $p < .05$, we retain the inequality for significance classification and use the numerical bound for eJAB. This can overstate evidence for the null, so the resulting contradiction classifications are provisional.

- *Section 2: 'More generally, the volume of prior probability ellipsoids around the MLE is $\mathcal{O}(n^{\frac{1-q}{2}})$.' So what? Why is this particular prior choice being made? Is this good or bad? Without appropriate justification, the proposed prior and BF are (for $q > 1$) ad-hoc.*

The empirically derived prior distribution underlying eJAB is

$$\boldsymbol{\theta} \sim \boldsymbol{\mathcal{N}}_q\left(\hat{\boldsymbol{\theta}},\ n^{\frac{1}{q}} \cdot \mathbf{J}_n^{-1}(\hat{\boldsymbol{\theta}})\right), \tag{8}$$

where $\mathbf{J}_n = \mathcal{O}(n)$ denotes the $q \times q$ information matrix. The volume of a fixed-probability ellipsoid is proportional to $|n^{\frac{1}{q}} \cdot \mathbf{J}_n^{-1}|^{\frac{1}{2}} = \mathcal{O}(n^{\frac{1-q}{2}})$. More generally, suppose that $\boldsymbol{\theta} \sim \boldsymbol{\mathcal{N}}_q\left(\hat{\boldsymbol{\theta}},\ a_n \cdot \mathbf{J}_n^{-1}(\hat{\boldsymbol{\theta}})\right)$, with $a_n > 1$. This yields $eJAB^*_{01} = a_n^{\frac{q}{2}} \exp\left\{-\frac{1}{2}\left(1 - \frac{1}{a_n}\right) Q_{\chi^2_q}(1-p)\right\}$ under the same Savage–Dickey approximation. Setting $a_n = n^{\frac{1}{q}}$ yields $\sqrt{n}$ growth in the evidence when $\mathcal{H}_0$ is true, so that under eJAB the rate of evidence accumulation towards $\mathcal{H}_0$ is independent of the dimension q. When $q = 1$, Eq. 8 reduces to the unit-information prior used in the original JAB (Aust et al., 2026; Wagenmakers, 2022). When $q > 1$ (F- and chi-squared tests) and effect sizes are small, it avoids the overly conservative behavior of the BIC approximation to the Bayes factor in favor of the null hypothesis (as detailed in the `R` Markdown files in the GitHub Simulations directory), while retaining model-selection consistency under R1 and R2.

Equally Weighted Decision Error Profiles: For each effect size θ and sample size n, we compare a simple null (effect zero) with a simple alternative (effect θ). Under 0–1 loss, the equally weighted decision error (EWDE) for decision rule d is

$$\mathrm{EWDE}_d(\theta, n) = \frac{\alpha_d(\theta, n) + \beta_d(\theta, n)}{2},$$

where α_d is the type I error probability and β_d is the type II error probability at the specified alternative. We define $\Delta E = \mathrm{EWDE}_{\mathrm{comparison}} - \mathrm{EWDE}_{\mathrm{eJAB}}$ and plot the normalized EWDE comparison,

$$\mathrm{EWDE}_{\mathrm{norm}}(\theta, n) = \frac{\Delta E(\theta, n)}{M}, \qquad M = \max_{(\theta,n)\in\mathcal{G}} |\Delta E(\theta, n)|,$$

where $\mathcal{G}$ is the displayed grid for each panel. Positive values favor eJAB; negative values favor the comparison rule.

Figure 2 of the revised manuscript now displays $\mathrm{EWDE}_{\mathrm{norm}}(\theta, n)$ and compares the performance of eJAB with: (A) TSBF where the effect size is estimated; (B) TSBF where the true effect size is assumed known; (C) $p \leq .005$; (D) the BIC approximation to the Bayes factor. The BIC comparison uses $BF_{01}^{\mathrm{BIC}} = n^{q/2}\exp(-\Lambda/2)$, where Λ is the likelihood-ratio statistic, for the six parametric-likelihood tests with available comparisons; Cox regression and the three rank tests are excluded.

The standalone $p \leq .05$ comparison is provided in Supplementary Figure S3 (Appendix B); both NHST comparisons are also shown below. The manuscript abstract and discussion now identify the observed regimes. For Pearson chi-squared independence tests in 2×2 and 3×3 tables, eJAB generally improves on estimated-ω TSBF outside small-sample or weak-signal regimes; small 3×3 samples can favor TSBF even at stronger effects. The pattern broadly reverses for 5×5 tables. With true ω, eJAB generally matches or underperforms TSBF, with differences that are not always small near weak signals.

Against BIC, single-parameter comparisons are close. The largest eJAB gains occur at weaker effects in ANOVA, repeated-measures ANOVA and chi-squared independence tests; at stronger effects, BIC can gain through lower type I error once both rules have high power. Against the NHST significance rule $p \leq .005$, eJAB gains in bands of weaker effects that move toward zero as sample size increases. NHST can have lower EWDE at stronger effects in smaller samples, particularly for multiparameter tests. At larger sample sizes, differences at stronger effects are small, while weak-signal bands can favor NHST. These are descriptions of the existing grids, not claims of uniform dominance.

A subset of the $\mathrm{EWDE}_{\mathrm{norm}}(\theta, n)$ plots making comparisons to TSBF, $p \leq .005$ and BIC are included as Figure 2 in the revision; the $p \leq .05$ comparison appears separately as Supplementary Figure S3.

- *Section 3: A table with the different designs (the ten tests), with an additional column for the degree of freedom, would be helpful. Then the second paragraph in Section 3.1 can be shortened.*

  We have added degrees of freedom to the panel facets and removed the second paragraph.

- *Section 3: 'effective sample size for survival analysis': A reference to Volinsky and Raftery (2000) might be useful. Also, it is not clear to me how the survival and censoring times were sampled.*

  We have added a reference to Volinsky and Raftery (2000) for the event-count definition of effective sample size. The survival designs discussed here are simulation studies. Survival analyses are excluded from the revised CTG application when the required event counts are unavailable, as documented in Appendix C.

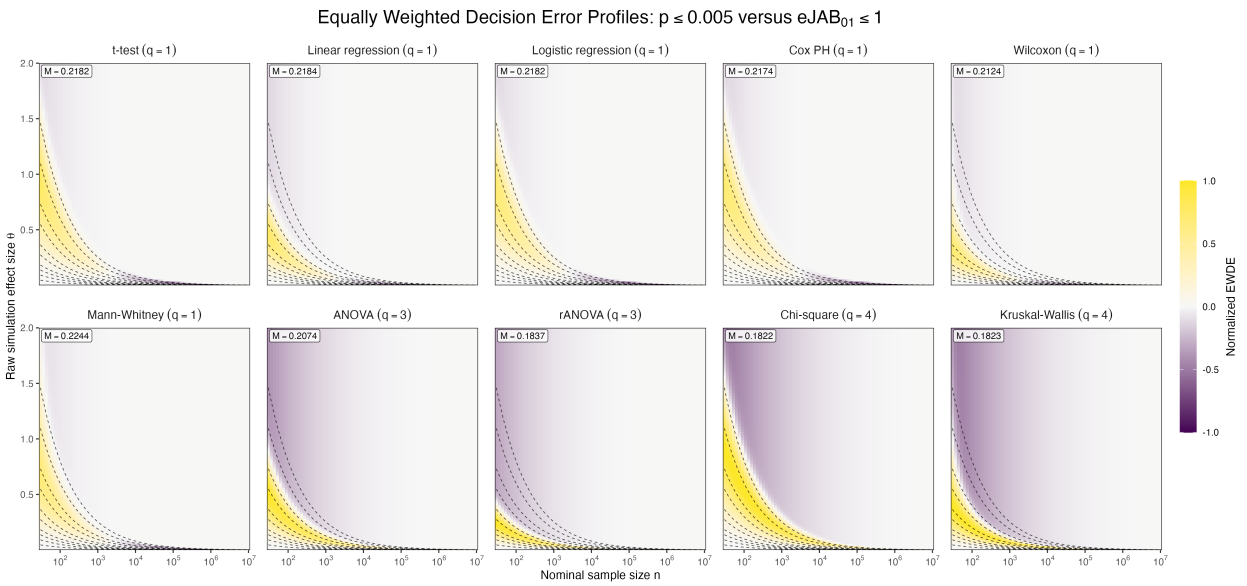


Figure 4: Equally Weighted Decision Error Profiles comparing eJAB with $p \leq .005$ across ten tests. The eJAB decision threshold is $BF_{01} \leq 1$. Colors show normalized EWDE; positive values favor eJAB and negative values favor the NHST significance rule.

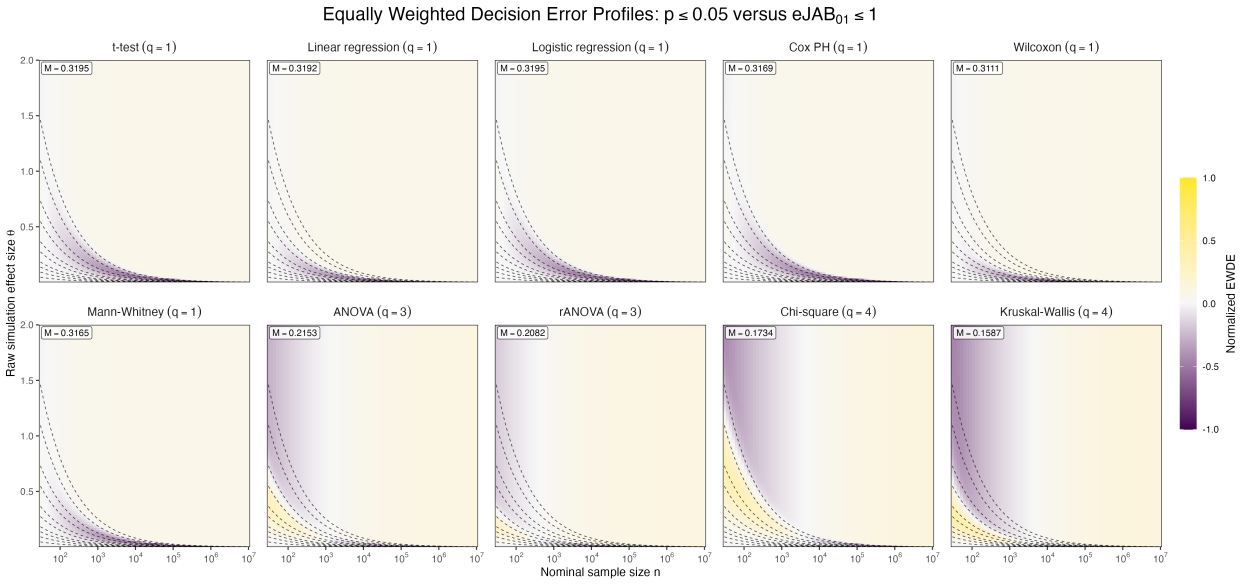


Figure 5: Equally Weighted Decision Error Profiles comparing eJAB with $p \leq .05$ across ten tests. The eJAB decision threshold is $BF_{01} \leq 1$. Colors show normalized EWDE; positive values favor eJAB and negative values favor the NHST significance rule.

In the simulations, survival times and censoring times are generated independently for each subject. The treatment indicator is binary, with approximately equal numbers of subjects assigned to the two groups and the assignments randomly permuted. Conditional on treatment status x_i, the survival time is sampled from an exponential distribution, $T_i \sim \text{Exp}(\lambda e^{bx_i})$ for $i = 1, \cdots, n$, so that the baseline hazard is λ and the treatment hazard ratio is e^b.

Right-censoring is imposed in one of two ways. In the simulations using random censoring, the censoring

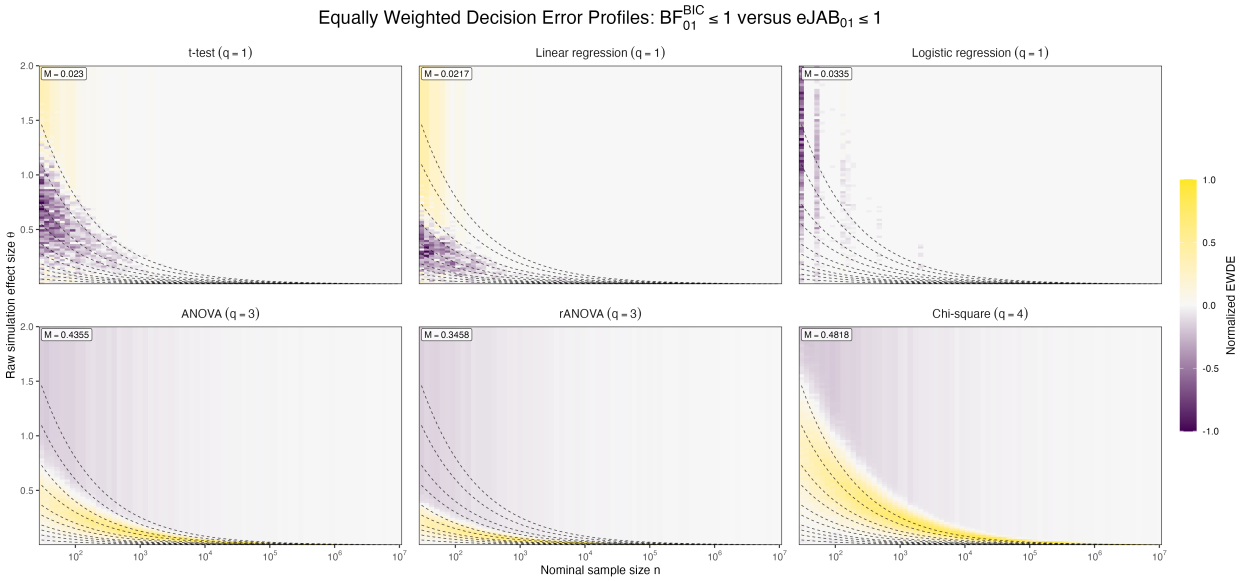


Figure 6: Equally Weighted Decision Error Profiles comparing eJAB with the BIC Bayes-factor rule across six parametric-likelihood tests: the two-sample t-test, linear regression, logistic regression, ANOVA, repeated-measures ANOVA, and chi-squared tests. Both Bayes-factor decision thresholds are $BF_{01} \leq 1$. Colors show normalized EWDE; positive values favor eJAB and negative values favor BIC. Each panel reports its normalization constant M, the maximum absolute EWDE difference over the displayed grid.

time is sampled independently as $C_i \sim \text{Exp}(\lambda_c)$ independently of both the survival time and treatment assignment, corresponding to non-informative right censoring. Alternatively, when a maximum follow-up time is specified, all subjects are administratively censored at the fixed time $C_i = T_{\max}$. The observed follow-up time is $Y_i = \min(T_i, C_i)$, with event indicator $\delta_i = \mathbf{1}_{T_i \leq C_i}$. Thus, subjects are sampled independently, and the censoring mechanism is non-informative by construction.

We used the number of observed events, $n^* = \sum_{i=1}^{n} \delta_i$, as the effective sample size in the eJAB formula. The simulation code and accompanying description are provided in `coxphSim.Rmd` in the GitHub repository.

- *'In particular, we tested the harmonic mean of group sizes for item 1, as in JAB (Wagenmakers, 2022), but do not demonstrate it here': I assume this is intended to investigate unequal sample sizes in the two groups? Should this not then be half the harmonic mean?*

The two sample-size considerations were not intended to address unequal group sizes per se. In the simulations, we used unequal group sizes with a ratio of 1.2:1, and we considered two sample-size conventions for computing Bayes factors: (1) the total number of observations and (2) one-half of the harmonic mean of the two group sizes.

The former convention is commonly used in linear-regression formulations and in the BIC approximation. The latter is also the inverse-variance effective sample size, $1/n = 1/n_1 + 1/n_2$, and is used in the JAB and JZS Bayes factors (Morey and Rouder, 2026). In the `ttestSim.Rmd` supplementary file in the GitHub repository, we included both sample-size definitions when computing JAB and BIC.

The use of one-half of the harmonic mean arises from the standard error of the difference between two independent group means. Wagenmakers (2022, pp. 26–27) discusses this sample-size choice in detail. Later, on p. 34, the central component is the $\sqrt{n}$ term—or, more generally, the sample-size quantity that enters the standard error of the maximum likelihood estimate.

- *Section 3.1: 'For local alternatives ...': Local alternatives only appear here, but they should also be discussed much more broadly in the theory section, as they form the basis of Johnson's BFs based on test-statistics.*

  We agree that local alternatives deserve a more prominent treatment in the theoretical development. We have expanded the discussion accordingly, including their role in TSBF and their relationship to our regularity conditions. These points are addressed in detail in the first bullet point of our response above.

- *Section 3.2: 'We compared eJAB with test-statistic BFs': How did you choose the parameter alpha (in Johnson's notation)? Or did you simply use the 'MMLE under the alternative hypothesis'? If so, it is not surprising that 'the test-statistic BF is inherently bounded below by 1'; it is actually a minimum BF! The discrepancies between your BF and Johnson's proposal are easy to explain, you are comparing apples with oranges.*

  We did indeed use the MMLE under the alternative hypothesis. In the revised manuscript, we instead compare the methods across several values of the hyperparameter and avoid comparisons based on the MMLE. These points are also addressed in detail in the first bullet point of our response above.

- *Beginning of Section 4: 'Prior work has raised concerns about the strength of statistical evidence in clinical research. For example, Begley and Ellis (2012) reported replication for only 11% of landmark preclinical cancer findings.' It does not make much sense to quote a study on the reproducibility of preclinical cancer findings as an example of 'concerns about the strength of statistical evidence in clinical research'. Clinical and preclinical research are very different fields.*

  Thank you for pointing this out. We have removed the reference from the revised manuscript.

- *The description of the 'candidate type I errors' is very repetitive, as you are always using the phrase 'We flag as a candidate type I error the reported effect ...' In addition, the reader would like to see additional information on sample size n (and dimension q) of the different studies, also because the focus is on the Jeffreys–Lindley paradox, a phenomenon that depends intrinsically on sample size. I noted that the first study is a four-arm Phase III trial, but only one p-value is reported. Is this the p-value for the comparison of all four groups or the p-value for the pairwise comparison of 50 mg vs placebo? What sample size was used, and did you also analyze the results from the other pairwise comparisons with different doses? If so, how did you deal with the dependence between the comparisons, given that they all share the same control arm?*

  In the revised manuscript, we report effect size, sample size, trial phase, and the outcome and comparison for each example. The three examples have $n = 66$, 80, and 1,586, respectively, with $q = 1$ in each case. The original four-arm Phase III trial is excluded by the revised eligibility criteria.

  More generally, a single study may contribute multiple p-values. For the posterior summaries, we aggregate results within studies into four categories defined by significance and whether eJAB exceeds 1. We use a Dirichlet$(1, 1, 1, 1)$ prior for each study and weight posterior category probabilities by the number of eligible analyses before calculating the contradiction proportion among significant results. By assigning each study its own category probabilities, the model allows for study-level heterogeneity.

- *End of Section 4: The linear model needs better motivation. Figure 4 would be more informative if the same y-axis scale were used throughout and if the panels were ordered A, C, B, D (for easier comparison of unadjusted and adjusted results).*

  Thank you for this suggestion. We have combined these results with the clinical-trial summaries in manuscript Figure 3. The outcome panels now appear as Figure 3D–G, with a common y-axis scale and unadjusted and adjusted results adjacent within each outcome type: D and E show primary outcomes, and F and G show secondary outcomes. This preserves the requested ordering of the former Figure 4 panels A, C, B, D.

  The linear mixed models describe within-study variation in $\ln eJAB$ and $\ln p$, separately for primary and secondary outcomes. We fit a fixed intercept and study-specific random intercept, and plot the residuals re-centered by the estimated fixed intercept.

References

Frederik Aust, Samuel Pawel, and Eric-Jan Wagenmakers. Extracting Bayesian evidence from frequentist p-values. *arXiv preprint arXiv:2607.12132*, 2026. doi: https://doi.org/10.48550/arXiv.2607.12132.

C. Glenn Begley and Lee M. Ellis. Raise standards for preclinical cancer research. *Nature*, 483(7391):531–533, 2012. doi: https://doi.org/10.1038/483531a.

Colin F. Camerer, Anna Dreber, Eskil Forsell, Teck-Hua Ho, Jürgen Huber, Magnus Johannesson, Michael Kirchler, Johan Almenberg, Adam Altmejd, Taizan Chan, Emma Heikensten, Felix Holzmeister, Taisuke Imai, Siri Isaksson, Gideon Nave, Thomas Pfeiffer, Michael Razen, and Hang Wu. Evaluating replicability of laboratory experiments in economics. *Science*, 351 (6280):1433–1436, 2016. doi: https://doi.org/10.1126/science.aaf0918.

Abdelaati Daouia, Gilles Stupfler, and Antoine Usseglio-Carleve. Corrected inference about the extreme Expected Shortfall in the general Max-Domain of Attraction. *Information and Inference: A Journal of the IMA*, 14(3):iaaf023, 2025. doi: https://doi.org/10.1093/imaiai/iaaf023.

Saptati Datta, Riana Guha, Rachael Shudde, and Valen E. Johnson. Bayes factors functions based on test statistics and non-local moment prior densities. *Statistics & Probability Letters*, 219:110330, 2025a. doi: https://doi.org/10.1016/j.spl.2024.110330.

Saptati Datta, Riana Guha, Rachael Shudde, and Valen E. Johnson. On Bayes factor functions. *Bayesian Analysis*, 20(4): 1399–1427, 2025b. doi: https://doi.org/10.1214/25-BA1562.

Timothy M. Errington, Elizabeth Iorns, William Gunn, Fraser Elisabeth Tan, Joelle Lomax, and Brian A. Nosek. An open investigation of the reproducibility of cancer biology research. *eLife*, 3:e04333, 2014. doi: https://doi.org/10.7554/eLife.04333.

Timothy M. Errington, Maya Mathur, Courtney K. Soderberg, Alexandria Denis, Nicole Perfito, Elizabeth Iorns, and Brian A. Nosek. Investigating the replicability of preclinical cancer biology. *eLife*, 10:e71601, 2021. doi: https://doi.org/10.7554/eLife.71601.

Leonhard Held and Manuela Ott. On p-values and Bayes factors. *Annual Review of Statistics and Its Application*, 5:393–419, 2018. doi: https://doi.org/10.1146/annurev-statistics-031017-100307.

Valen E. Johnson. Bayes factors based on test statistics. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 67(5):689–701, 2005. doi: https://doi.org/10.1111/j.1467-9868.2005.00521.x.

Valen E. Johnson. Properties of Bayes factors based on test statistics. *Scandinavian Journal of Statistics*, 35(2):354–368, 2008. doi: https://doi.org/10.1111/j.1467-9469.2007.00576.x.

Valen E. Johnson. Revised standards for statistical evidence. *Proceedings of the National Academy of Sciences*, 110(48): 19313–19317, 2013. doi: https://doi.org/10.1073/pnas.1313476110.

Valen E. Johnson, Sandipan Pramanik, and Rachael Shudde. Bayes factor functions for reporting outcomes of hypothesis tests. *Proceedings of the National Academy of Sciences*, 120(8):e2217331120, 2023. doi: https://doi.org/10.1073/pnas.2217331120.

Richard D. Morey and Jeffrey N. Rouder. BayesFactor: Computation of Bayes factors for common designs, 2026. URL `https://github.com/richarddmorey/bayesfactor`. R package version 0.9.12-4.8.

Open Science Collaboration. Estimating the reproducibility of psychological science. *Science*, 349(6251):aac4716, 2015. doi: https://doi.org/10.1126/science.aac4716.

Thomas Sellke, María Jesús Bayarri, and James O. Berger. Calibration of ρ values for testing precise null hypotheses. *The American Statistician*, 55(1):62–71, 2001. doi: https://doi.org/10.1198/000313001300339950.

Jason Travers and Michael Krezmien. Racial disparities in autism identification in the united states during 2014. *Exceptional Children*, 84(4):403–419, 2018. doi: 10.1177/0014402918771337.

Alexander H. Tyner et al. Investigating the replicability of the social and behavioural sciences. *Nature*, 652:143–151, 2025. doi: https://doi.org/10.1038/s41586-025-10078-y.

Chris T. Volinsky and Adrian E. Raftery. Bayesian information criterion for censored survival models. *Biometrics*, 56(1):256–262, 2000. doi: https://doi.org/10.1111/j.0006-341x.2000.00256.x.

Eric-Jan Wagenmakers. Approximate objective Bayes factors from p-values and sample size: The $3p\sqrt{n}$ rule. *PsyArXiv*, 1:1–50, 2022. doi: https://doi.org/10.31234/osf.io/egydq.

Ying Yuan and Valen E. Johnson. Bayesian hypothesis tests using nonparametric statistics. *Statistica Sinica*, 18(3):1185–1200, 2008. URL `https://www.jstor.org/stable/24308537`.

Response to Reviewer 2

There is no doubt that this work will attract a lot of attention. Similar to the 2018 Benjamin et al. paper, the approach outlined in this paper is straightforward to implement, and, in my opinion, there is a good argument to be made that when $p < \alpha$ ('reject $\mathcal{H}_0$') while $BF_{01} > 1$ ('the data support $\mathcal{H}_0$'), or even < 1 but $> 1/3$ ('the data provide weak evidence against $\mathcal{H}_0$'), this casts substantial doubt on the frequentist conclusion – at any rate, such a conflict demands a more careful analysis. The fact that such conflicts now go entirely unnoticed seems unproductive. It may be that, at several points throughout the paper, its accessibility can be increased even more. Regardless, the work provides a fresh perspective on a crucial issue that continues to haunt the field. It is certain to lead to a useful debate, and, in fact, I would suggest that –if it is accepted– it will be the target of a few comments.

Thank you very much for this encouraging and thoughtful assessment. We appreciate the reviewer's recognition of the practical relevance and potential impact of the proposed approach. We also agree that conflicts between NHST and Bayes-factor conclusions warrant closer scrutiny rather than going unnoticed (e.g., Aust et al., 2026; Tyner et al., 2025).

In the revised manuscript, we have expanded the theoretical content to provide a stronger justification for the methodology. We also agree on the importance of making the methodology more accessible. To this end, we have developed an eJAB calculator, available at `https://flxp0x-puneet-velidi.shinyapps.io/bayesfactor/`, and implemented eJAB in a JASP module designed to identify potentially irreplicable results. A companion article describes the JASP module and presents three examples in which it is used to identify, from a collection of studies, those that are unlikely to replicate when replication outcomes are subsequently available.

The manuscript's Equally Weighted Decision Error Profiles compare eJAB with the NHST significance rule $p \leq .005$, TSBFs evaluated at true and estimated effects, and BIC approximations; Supplementary Figure S3 retains the $p \leq .05$ comparison. These profiles show where each rule has lower decision error across sample sizes, effect sizes, and dimensions.

Frederik Aust, Samuel Pawel, and Eric-Jan Wagenmakers. Extracting Bayesian evidence from frequentist p-values. *arXiv preprint arXiv:2607.12132*, 2026.

Alexander H. Tyner et al. Investigating the replicability of the social and behavioural sciences. *Nature*, 652: 143–151, 2025. doi: 10.1038/s41586-025-08258-3.