

Appendix A: Preprocessing Steps

CTG Download. We downloaded the CTG results for completed studies with posted results as a JSON export (all available fields) from <https://clinicaltrials.gov/search?aggFilters=results:with,status:com> on August 7th, 2025. During data cleaning, we excluded four records with sample-size errors identified while reviewing candidate type I errors at $\alpha = .01$.

Record Construction. We transformed each ClinicalTrials.gov JSON record into flat, analysis-level rows prior to p -value cleaning and method harmonization. For each study (`nctId`), we iterated over outcome measures with at least one analysis and, for every outcome–analysis pair, emitted one row per denominator count with units “Participants” and with `groupId` in the analysis’s declared `groupIds`. We carried forward study metadata (e.g., `briefTitle`, `studyType`), design-information fields (`interventionModel`, `observationalModel`, `timePerspective`, `allocation`), one-hot encoded phases (`EARLY_PHASE1`, `PHASE1`, `PHASE2`, `PHASE3`, `PHASE4`), and MeSH terms/IDs for conditions and interventions from the `derivedSection`. We also flagged whether the study included any references of type `RESULTS` or `DERIVED`.

Analysis Identifier. Each row was keyed by a stable, hash-based identifier,

$$\text{analysisId} = \text{MD5}(\text{nctId} + \text{JSON}(\text{outcome}; \text{sorted keys}) + \text{JSON}(\text{analysis}; \text{sorted keys})),$$

which yields reproducible IDs for unique outcome–analysis pairs within a study (independent of JSON key order). This `analysisId` serves as the unit of aggregation in downstream preprocessing (e.g., computing sample sizes and $eJAB$).

Schema. Fields include: `analysisId`, `nctId`, `briefTitle`, `outcomeType`, `outcomeTitle`, `groupDescription`, `units`, `groupId`, `value`, `pValue`, `statisticalMethod`; MeSH fields (`condition_mesh`, `condition_ids`, `intervention_mesh`, `intervention_ids`); `hasResultsOrDerived`; study-design fields; and phase indicators. The field `value` records the number of participants for the given outcome group/arm.

Standardization of Reported p -Values. Reported p -values often included textual qualifiers (e.g., “<0.05”). We standardized them as follows:

1. Flag the presence of a “<” qualifier: `has_less_than` $\leftarrow$ `grepl("<", pValue)`.
2. Extract the numeric component via regex:

$$\text{pValue} \leftarrow \text{str_extract}(\text{pValue}, "[0-9]*\.[0-9]+").$$

3. Drop rows with missing or blank numeric components.
4. Coerce `pValue` to numeric and retain only valid probabilities: $0 < \text{pValue} < 1$.
5. Exclude rows with zero sample size, which are incompatible with downstream model mappings.

Harmonization of Statistical Method Names. Free-text `statisticalMethod` labels were mapped to canonical test families using regular-expression rules (`str_detect`). Observations with no match were excluded. The exact mapping rules are:

Listing 1: Regex-based mapping of free-text methods to canonical families

```
statisticalMethod <- case_when(  
  # --- T-TESTS ---  
  str_detect(originalMethod, regex("\\bt[- ]?test", ignore_case = TRUE)) &  
    str_detect(originalMethod, regex("one[- ]?sample", ignore_case = TRUE)) ~ "1-ttest",  
  
  str_detect(originalMethod, regex("\\bt[- ]?test", ignore_case = TRUE)) &  
    str_detect(originalMethod, regex("paired", ignore_case = TRUE)) ~ "1-ttest",  
  
  str_detect(originalMethod, regex("\\bt[- ]?test", ignore_case = TRUE)) ~ "2-ttest",  
  
  # --- CONDITIONAL LOGISTIC REGRESSION ---  
  str_detect(originalMethod, regex("logistic", ignore_case = TRUE)) &  
    str_detect(originalMethod, regex("regression", ignore_case = TRUE)) &  
    str_detect(originalMethod, regex("conditional", ignore_case = TRUE)) ~ "clogit",  
  
  # --- LOGISTIC REGRESSION ---
```

```

str_detect(originalMethod, regex("logistic", ignore_case = TRUE)) &
  str_detect(originalMethod, regex("regression", ignore_case = TRUE)) ~ "glm",

# --- COX ---
str_detect(originalMethod, regex("cox", ignore_case = TRUE)) &
  !str_detect(originalMethod, regex("wilcoxon|mantel|signed|rank", ignore_case = TRUE)) ~ "cox",

# --- LOGRANK / MANTEL-COX ---
str_detect(originalMethod, regex("log ?rank|mantel[- ]?cox", ignore_case = TRUE)) &
  !str_detect(originalMethod, regex("wilcoxon|signed", ignore_case = TRUE)) ~ "logrank",

# --- MANN-WHITNEY ---
str_detect(originalMethod, regex("mann[- ]?whitney", ignore_case = TRUE)) |
  (str_detect(originalMethod, regex("wilcoxon", ignore_case = TRUE)) &
    str_detect(originalMethod, regex("rank", ignore_case = TRUE)) &
    !str_detect(originalMethod, regex("signed", ignore_case = TRUE))) ~ "Utest",

# --- WILCOXON SIGNED-RANK ---
str_detect(originalMethod, regex("wilcoxon", ignore_case = TRUE)) &
  str_detect(originalMethod, regex("signed", ignore_case = TRUE)) &
  str_detect(originalMethod, regex("rank", ignore_case = TRUE)) ~ "wilcox",

# --- KRUSKAL-WALLIS ---
str_detect(originalMethod, regex("kruskal|wallis", ignore_case = TRUE)) ~ "Htest",

# --- CHI-SQUARE ---
str_detect(originalMethod, regex("chi[- ]?square|χ2|chi2|chisq", ignore_case = TRUE)) ~ "chisq",

# --- REPEATED MEASURES ---
str_detect(originalMethod, regex("repeated[- ]?measures?", ignore_case = TRUE)) ~ "rANOVA",

# --- LINEAR REGRESSION ---
str_detect(originalMethod, regex("linear", ignore_case = TRUE)) &
  str_detect(originalMethod, regex("regression", ignore_case = TRUE)) ~ "lm",

# --- ANOVA ---
str_detect(originalMethod, regex("anova|analysis of variance", ignore_case = TRUE)) ~ "ANOVA",

TRUE ~ NA_character_
)

```

Adjustment for One-Sided p -Values. When free text indicated a one-sided analysis (patterns: one-sided, 1-sided), we converted to a two-sided p -value by doubling and truncating at 1:

$$pValue \leftarrow \min\{2 \cdot pValue, 1\}.$$

This rule was applied only when `str_detect` flagged these patterns in `originalMethod`.

analysisId Sample Size QC. Certain tests should yield exactly one row per `analysisId` (e.g., one-sample t -test, Wilcoxon signed-rank test). To enforce consistency, we dropped analyses in which such tests appeared with multiple rows under the same `analysisId`.

Aggregation to the Test Level and Parameter Construction. We aggregated records at the `analysisId` (test) level. Fields invariant within a test were taken from the first row—e.g., `nctId`, `pValue`, `statisticalMethod`, `studyType`, phase indicators (EARLY_PHASE1, PHASE1, PHASE2, PHASE3, PHASE4), `outcomeTitle`, `outcomeType`, `groupDescription`, `conditionIds`, `briefTitle`, and `has_less_than`.

We defined the effective sample size n per test as:

$$n = \begin{cases} \text{first}(\text{value}), & \text{one-sample } t\text{-test or Wilcoxon signed-rank;} \\ \max \text{ value}, & \text{conditional logistic regression;} \\ 0.5 \sum \text{ value}, & \text{Cox PH or log-rank;} \\ \sum \text{ value}, & \text{otherwise.} \end{cases}$$

We also set method-specific parameters required by *eJAB*:

$$I = \begin{cases} n_{\text{rows within test}}, & \text{ANOVA, Kruskal-Wallis, repeated measures;} \\ \text{NA}, & \text{otherwise,} \end{cases}$$

$$(R, C) = \begin{cases} (n_{\text{rows within test}}, 2), & \text{chi-squared test;} \\ (\text{NA}, \text{NA}), & \text{otherwise.} \end{cases}$$

Mapping to JAB Models and Computation of eJAB. We mapped canonical `statisticalMethod` labels to the JAB01 model strings required for $eJAB_{01}$. Mapping was performed only for tests included in our analysis (i.e., excluding chi-squared, Cox PH, log-rank, and repeated-measures designs):

Two-sample t -test, paired t -test, one-sample t -test $\rightarrow$ "t-test",
 Linear regression $\rightarrow$ "linear_regression",
 Logistic regression, conditional logistic regression $\rightarrow$ "logistic_regression",
 ANOVA $\rightarrow$ "anova", Kruskal–Wallis $\rightarrow$ "kruskal_wallis",
 Mann–Whitney $\rightarrow$ "mann_whitney", Wilcoxon $\rightarrow$ "wilcoxon".

Given $(n, pValue, model, R, C, I)$ for each test, we computed $eJAB_{01}$ via

$$JAB = JAB01(n, pValue, model; R, C, I).$$

Exclusions and Final Dataset. We excluded rows with missing `model`, missing JAB or `pValue`, and tests with $n = 1$. The manuscript analysis then applied the criteria in Appendix C to obtain the same screened set used throughout: 16,563 eligible primary- and secondary-outcome analyses from 1,790 exact Early Phase I, Phase I, and Phase II trials.

Appendix B: Supplementary Figures




Figure S2: Comparison of the natural logarithmic $eJAB_{10}$ with the `contingencyTableBF` function from the ‘*BayesFactor*’ R package at three total counts for chi-squared tests of independence under two sampling designs: multinomial (top) and product-multinomial (bottom). Colors represent the true effect size. The dashed lines indicate $eJAB_{10}$ of $1/3$ and 3 , representing moderate evidence against and for an effect, respectively.




Figure S4: Scatterplots of $\ln eJAB_{01}$ versus p -value by MeSH condition (part 1 of 3) in the screened primary- and secondary-outcome set, with inset zooms for $p \leq .05$. Shaded bands denote evidence categories.




Figure S5: Scatterplots of $\ln eJAB_{01}$ versus p -value by MeSH condition (part 2 of 3) in the screened primary- and secondary-outcome set, with inset zooms for $p \leq .05$. Shaded bands denote evidence categories.




Figure S6: Scatterplots of $\ln eJAB_{01}$ versus p -value by MeSH condition (part 3 of 3) in the screened primary- and secondary-outcome set, with inset zooms for $p \leq .05$. Shaded bands denote evidence categories.




Figure S7: Residuals versus fitted values from random-intercept models for $\ln eJAB_{01}$ (left) and $\ln p$ (right) in the screened primary- and secondary-outcome set. The dashed horizontal line marks zero residuals and the solid curve is a LOESS smooth. Nearly parallel slanted bands reflect heaping at common reported p thresholds.




Figure S8: Distributions of p -values (A; vertical dashed line at .05) and sample size (B) by exact Early Phase I, Phase I, and Phase II classification in the screened primary- and secondary-outcome set.

Appendix C: Detailed ClinicalTrials.gov Inclusion and Exclusion Criteria

C.1 Scope, archived source, and units

The source was the ClinicalTrials.gov (CTG) JSON results export for completed studies with posted results, downloaded on August 7, 2025. The extraction emitted one row for every participant-denominator group attached to an outcome-analysis pair. A deterministic MD5 hash of the CTG identifier, the complete outcome object, and the complete analysis object (with sorted JSON keys) defined `analysisId`. Thus, `analysisId`, rather than an individual denominator row or a CTG study, was the unit of statistical analysis. Rows sharing an `analysisId` were group-denominator components of the same test and were collapsed during sample-size construction. Statistical analyses were nested within CTG studies identified by `nctId`.

The manuscript analysis estimates the frequency of disagreement between direction-specific statistical significance and eJAB evidence among eligible primary- and secondary-outcome analyses from exact Early Phase I, exact Phase I, and exact Phase II interventional trials. The three phase classifications are retained as separate groups. These criteria define the screened set used throughout the manuscript and supplement.

C.2 Inclusion criteria

An analysis was eligible only when all of the following criteria were met:

1. The record occurred in the archived CTG results extraction and had a deterministic `analysisId` and `nctId`.
2. CTG classified the study type as `INTERVENTIONAL`.
3. CTG classified the study as *exactly* Early Phase I, *exactly* Phase I, or *exactly* Phase II. Exact Early Phase I required `EARLY_PHASE1=1` and `PHASE1=PHASE2=PHASE3=PHASE4=0`. Exact Phase I and exact Phase II were defined analogously using their respective indicator.
4. CTG labeled the outcome as `PRIMARY` or `SECONDARY`.
5. The reported p -value contained a numeric value strictly between 0 and 1 and was not right-censored (for example, $p > .05$).
6. The free-text statistical method mapped deterministically to one of the eJAB-computable classes in Table 2.
7. All participant denominators required for the mapped test were finite and positive; the derived sample size satisfied $n > 1$; and the test dimension satisfied $q \geq 1$.
8. The resulting direction-appropriate eJAB was finite and positive.

Statistical significance was *not* an eligibility criterion. Both significant and nonsignificant eligible results were retained for the four-category study-level model.

C.3 Exclusion criteria

The following were excluded, in the sequential order shown in Table 1:

- Four prespecified malformed analysis identifiers previously found to contain denominator errors.
- Observational and expanded-access records, and any record not labeled interventional.
- Phase I/Phase II, Phase II/Phase III, Phase III, Phase IV, not-applicable, or missing phase classifications. Combined phases were not recoded into any of the three early-phase groups.
- Outcomes other than primary or secondary, baseline measures, participant-flow analyses, and adverse-event analyses.

- Missing, blank, nonnumeric, zero, one, or otherwise invalid p -values. Right-censored values were excluded because their numerical bound does not identify the value required for eJAB. Left-censored values were handled as described in Section C.5.
- Methods outside the prespecified mappings, including ANCOVA, mixed-effects or repeated-measures models, Fisher exact tests, chi-squared tests without full table dimensions, Cox or log-rank analyses without the number of observed events, Bayesian analyses, and other free-text methods that could not be mapped without assumptions not supported by the CTG record.
- Analyses with missing, zero, or ambiguous participant denominators; one-sample or signed-rank analyses represented by more than one denominator row; omnibus analyses represented by fewer than two groups; $n \leq 1$; or $q < 1$.
- A nonfinite or nonpositive eJAB value after all preceding steps.

Table 1: Sequential inclusion and exclusion flow. Counts are analysis identifiers except in the last column. Each exclusion count is conditional on all previous steps.

Screening stage	Excluded	Retained	Studies
Unique analyses in the archived extraction	0	279,927	22,454
Remove four malformed analysis identifiers	4	279,923	22,453
Restrict to interventional studies	11,893	268,030	21,500
Restrict to exact Early Phase I, Phase I, or Phase II	175,191	92,839	6,302
Restrict to primary and secondary outcomes	1,806	91,033	6,277
Require an admissible p -value and exclude right-censoring	26,665	64,368	4,461
Require a prespecified eJAB method mapping	47,214	17,154	1,877
Require valid denominators, row structure, n , and q	591	16,563	1,790

C.4 Method mapping, dimension, and sample size

Free-text method labels were normalized using case-insensitive regular expressions. More specific patterns were evaluated before broader patterns (for example, conditional logistic regression before logistic regression, and one-sample or paired t -tests before other t -tests). Table 2 gives the complete set used for the early-phase reanalysis.

Table 2: Eligible statistical method classes and derived quantities. Here G is the number of attached denominator groups.

Mapped class	Sample-size rule	Dimension q
One-sample or paired t -test	One attached denominator; n is that value	1
Two-sample t -test	Sum attached group denominators	1
Linear regression	Sum attached denominators	1
Logistic regression	Sum attached denominators	1
Conditional logistic regression	Maximum attached denominator	1
One-way ANOVA	Sum attached group denominators; require $G \geq 2$	$G - 1$
Wilcoxon signed-rank test	One attached denominator; n is that value	1
Mann–Whitney test	Sum attached group denominators	1
Kruskal–Wallis test	Sum attached group denominators; require $G \geq 2$	$G - 1$

Participant counts containing thousands separators were parsed as numbers before aggregation. A missing or nonpositive attached denominator made the full analysis ineligible; missing values were not silently removed from a sum. The archived extraction’s repeated `analysisId` rows encode attached denominator groups and were therefore aggregated, not counted as duplicate hypothesis tests. The hash construction produced one final record per outcome–analysis pair; no additional duplicate analysis identifiers remained after aggregation.

C.5 Reported inequalities, test direction, and Wagenmakers’s one-sided Bayes factor

The numeric component of a reported p -value was extracted reproducibly. For a left-censored value such as $p < .05$, the bound .05 was used in the eJAB calculation. This is conservative for evidence against the null because the unknown underlying value is smaller. The inequality itself was retained: $p < .05$ met a .05 significance rule even though the stored numerical bound equaled .05. Values reported with “>” were excluded because a lower bound does not identify the tail probability needed for eJAB. The final eligible set contained 2,714 left-censored values (30 in Early Phase I, 429 in Phase I, and 2,255 in Phase II).

Direction was determined from the reported `statisticalMethod`. The patterns `one-sided`, `1-sided`, `one-tailed`, and `1-tailed` identified a one-sided test. A directional correction was used only for scalar tests ($q = 1$). An explicit two-sided label or the absence of an explicit one-sided method label was classified as two-sided/unrestricted. Mentions of one-sided α found only in sample-size or power prose were not allowed to reclassify the reported test. This rule identified 427 one-sided analyses and 16,136 two-sided/unrestricted analyses.

For a one-sided reported value p_1 , the corresponding two-sided value was first reconstructed as

$$p = 2 \min\{p_1, 1 - p_1\}.$$

The ordinary two-sided eJAB was then computed as

$$eJAB_{01} = \sqrt{n} \exp\left[-\frac{1}{2} \left(1 - n^{-1/q}\right) Q_{\chi^2_q}(1 - p)\right].$$

Following the encompassing-prior/order-restriction approximation of Wagenmakers (2022, pp. 20–22), the directional factor was

$$eJAB_{0+} = \frac{eJAB_{01}}{2(1 - p_1)}.$$

The denominator is the approximate posterior-to-prior mass correction for the predicted half of a symmetric alternative: $\Pr(\theta > \theta_0 \mid \mathcal{H}_1) = 1/2$ a priori and approximately $1 - p_1$ a posteriori for moderate samples and a relatively vague prior. Thus the procedure does not merely substitute a one-sided p -value into a chi-squared formula.

Statistical significance was defined as $p_1 < .025$ for one-sided analyses and $p < .05$ for two-sided/unrestricted analyses. A significant result was called a *Bayes/NHST contradiction* when the applicable $eJAB_{01}$ or $eJAB_{0+}$ exceeded 1. This descriptive term does not assert that the result is a type I error.

C.6 Final analytic set and reproducibility

The final set comprised 16,563 primary- and secondary-outcome analyses from 1,790 trials: 270 analyses from 47 exact Early Phase I trials, 2,098 analyses from 280 exact Phase I trials, and 14,195 analyses from 1,463 exact Phase II trials. Of these, 90 Early Phase I, 918 Phase I, and 5,401 Phase II analyses met their direction-specific significance threshold. Five Early Phase I results (5.56% of significant results), 25 Phase I results (2.72%), and 431 Phase II results (7.98%) were Bayes/NHST contradictions at the applicable direction-specific $eJAB > 1$ threshold.

The complete executable pipeline is `response_to_reviewers/analysis/early_phase_analysis.R`. It writes the sequential flow table, the eligible analysis-level data set, phase and test summaries, posterior draws, and the clustered mixed-model sensitivity result. The four manual exclusions, method regular expressions, direction rules, significance thresholds, denominator rules, and eJAB calculations are encoded directly in that script.

The July 2026 protocol was archived after the August 2025 CTG snapshot had already been downloaded and inspected. Accordingly, this is described as a protocol-aligned retrospective reanalysis, not as a prospective confirmatory study or an independent replication. The phase-specific posterior proportions are descriptive, and no confirmatory comparison among phases is made.

Appendix D: Proof of Model-Selection Consistency

Lemma 1. For $x \rightarrow 1^-$,

$$0 < f(x) = -2\ln(1-x) - 2\ln(-\ln(1-x)) < Q_{\chi_q^2}(x).$$

The proof is provided in the GitHub repository (see files within ‘Theorem’).

PROOF OF THEOREM 2.2. Assume $\mathcal{H}_0$ is true. Let $\delta \in (0, 1)$ and $K > 0$ be arbitrary. Note that **R1** and $W_n = Q_{\chi_q^2}(1-p) \xrightarrow{D} \chi_q^2$ by the continuous mapping theorem.

For $n \geq N^* = \lceil K^2 \exp\{Q_{\chi_q^2}(1-\delta)\} \rceil$, we have

$$\begin{aligned} \mathbb{P}(eJAB_{01} > K) &= \mathbb{P}\left(\sqrt{n} \exp\left\{-\frac{1}{2} \frac{n^{1/q} - 1}{n^{1/q}} W_n\right\} > K\right) \\ &> \mathbb{P}\left(\sqrt{n} \exp\left\{-\frac{1}{2} W_n\right\} > K\right) \\ &= \mathbb{P}\left(W_n < -2 \ln \frac{K}{\sqrt{n}}\right) \\ &\geq \mathbb{P}\left(W_n < -2 \ln \frac{K}{\sqrt{N^*}}\right) \\ &\geq \mathbb{P}\left(W_n < -2 \ln \frac{K}{\sqrt{K^2 \exp\{Q_{\chi_q^2}(1-\delta)\}}}\right) \\ &= \mathbb{P}\left(W_n < Q_{\chi_q^2}(1-\delta)\right) \rightarrow 1 - \delta \text{ as } n \rightarrow \infty. \end{aligned}$$

Thus, $eJAB_{01} \xrightarrow{P} \infty$ as $n \rightarrow \infty$ under $\mathcal{H}_0$.

Assume $\mathcal{H}_1$ is true. Let $\delta \in (0, 1)$ and $\epsilon > 0$ be arbitrary. Lemma 1 gives $0 < f(1-p) < Q_{\chi_q^2}(1-p)$ for $p \rightarrow 0$.

$$\begin{aligned} &\mathbb{P}(eJAB_{01} < \epsilon) \\ &= \mathbb{P}\left(\sqrt{n} \exp\left\{-\frac{1}{2} \frac{n^{1/q} - 1}{n^{1/q}} Q_{\chi_q^2}(1-p)\right\} < \epsilon\right) \\ &> \mathbb{P}\left(\sqrt{n} \exp\left\{-\frac{1}{2} \frac{n^{1/q} - 1}{n^{1/q}} f(1-p)\right\} < \epsilon\right) \\ &= \mathbb{P}\left(\sqrt{n} (p \cdot (-\ln p))^{\frac{n^{1/q}-1}{n^{1/q}}} < \epsilon\right) > 1 - \delta. \end{aligned}$$

The final inequality is true for n sufficiently large by **R2** and the continuous mapping theorem with $1 - \frac{1}{n^{1/q}} \rightarrow 1$.

Thus, $eJAB_{01} \xrightarrow{P} 0$ as $n \rightarrow \infty$ under $\mathcal{H}_1$. Together, $eJAB$ is model-selection consistent. □